■ Physics Study Buddy

Agentic AI Capstone Project — 2026

Dr. Kanthi Kiran Sirra | Sr. AI Engineer | Agentic AI Course 2026

Field	Detail
Domain	Study Buddy — B.Tech Physics
User	Engineering students needing concept help at any hour
Tool Used	Calculator — for numerical physics problems (F=ma, KE=1/2 mv^2)
LLM	LLaMA 3.3-70B via Groq API
Vector DB	ChromaDB (in-memory, 10 documents)
Framework	LangGraph StateGraph + MemorySaver
UI	Streamlit (web browser interface)

1. Problem Statement

B.Tech Physics students often struggle with concepts and numerical problems outside classroom hours when no teacher is available. Existing general-purpose chatbots frequently hallucinate formulas and present incorrect physics derivations with high confidence, making them unreliable for academic use. Students need a 24/7 assistant that is strictly grounded in their syllabus, never invents answers, and can also perform calculations accurately.

Key requirements identified:

- Answer questions faithfully from a verified 10-topic physics knowledge base.
- Perform arithmetic calculations for numerical problems using a safe calculator tool.
- Remember student name and context across multiple turns within a session.
- Clearly admit when a question is outside its syllabus — no hallucination under any circumstance.

2. Solution and Features

Physics Study Buddy is a Retrieval-Augmented Generation (RAG) agent built on a LangGraph StateGraph with eight specialised nodes. It combines semantic search over a curated physics knowledge base with a self-reflection evaluation loop that ensures answer quality before returning a response to the student.

2.1 Agent Architecture

The agent follows a fixed 8-node pipeline. Each node reads from and writes to a shared CapstoneState TypedDict, ensuring all state transitions are explicit and testable.

Node	Role
memory_node	Appends question to history (sliding window of 6), extracts student name
router_node	LLM classifies question as retrieve / tool / memory_only
retrieval_node	Embeds query, fetches top-3 ChromaDB chunks with topic labels
skip_node	Returns empty context for greetings and small-talk
tool_node	Extracts math expression and evaluates using safe Python eval
answer_node	Generates grounded answer using system prompt + context + history
eval_node	Scores faithfulness 0.0-1.0; triggers retry if below 0.7 (max 2 retries)
save_node	Persists final answer to conversation history, then END

2.2 Key Features

Strict Grounding: The system prompt forbids the LLM from using general knowledge. Every answer must come only from retrieved context or the calculator result.

Self-Reflection Eval: After every answer, eval_node scores faithfulness using the LLM as a judge. Scores below 0.7 trigger a retry with an escalated instruction. Maximum 2 retries prevent infinite loops.

Multi-Turn Memory: MemorySaver persists full graph state per thread_id across multiple invoke() calls. A sliding window of 6 messages prevents token overflow on the

Groq free tier.

Calculator Tool: A safe Python eval sandbox handles numerical physics problems. The LLM extracts the expression; the sandbox computes it. Tools never raise exceptions — errors return as strings.

Honest Refusal: Out-of-scope and adversarial questions are cleanly rejected with a standard message directing students to their professor.

2.3 Technology Stack

Component	Technology	Purpose
LLM	LLaMA 3.3-70B (Groq)	Routing, answering, evaluation
Embeddings	all-MiniLM-L6-v2	Semantic similarity for retrieval
Vector DB	ChromaDB in-memory	Stores and queries KB documents
Orchestration	LangGraph StateGraph	Node pipeline and conditional routing
Memory	LangGraph MemorySaver	Multi-turn session persistence
UI	Streamlit	Web chat interface
Evaluation	RAGAS	Faithfulness, relevancy, precision metrics

3. Application Screenshots

3.1 Home Screen and Concept Q&A

The Streamlit interface shows the sidebar listing all 10 physics topics, the chat area, and the agent's grounded answer to 'What is Newton's second law?' — retrieved directly from the knowledge base with the formula F = ma.

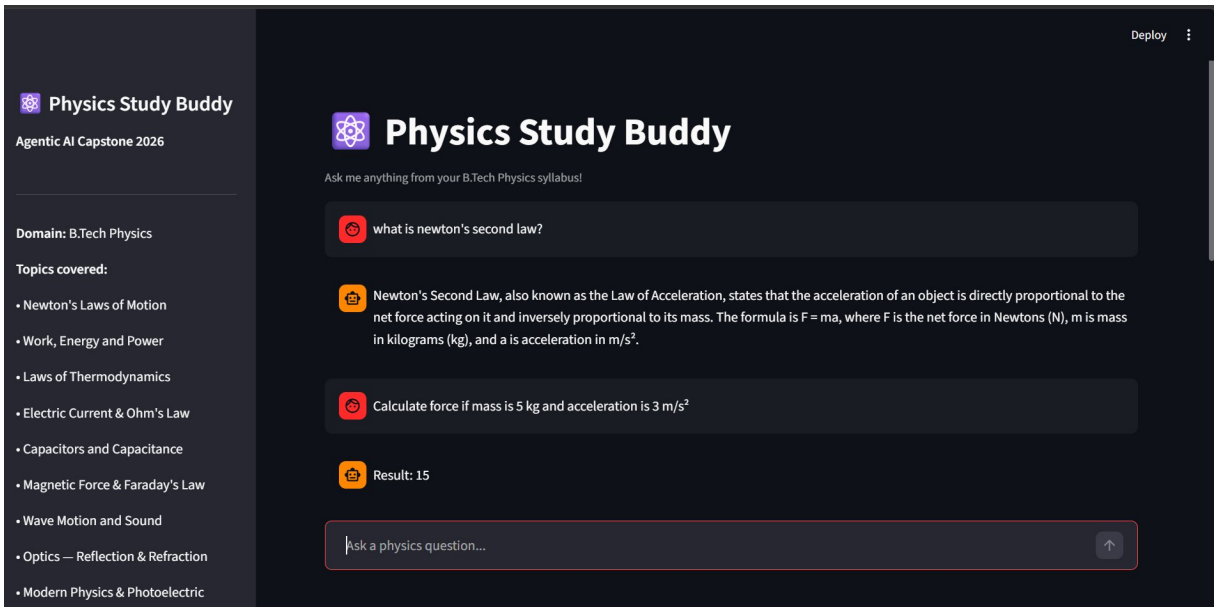

Figure 1: Home screen with concept Q&A; response

3.2 Calculator Tool in Action

When the student asks a numerical question, the router sends it to the tool node. The agent extracts the expression '5 * 3' and returns the correct result: 15 N.

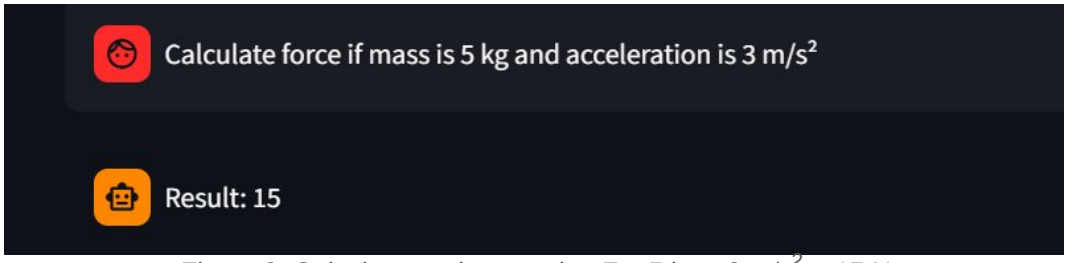

Figure 2: Calculator tool computing F = 5 kg x 3 m/s^2 = 15 N

3.3 Out-of-Scope Refusal

When asked 'What is GDP of India?' — a question completely outside the physics syllabus — the agent correctly refuses and directs the student to their professor, with no hallucination.

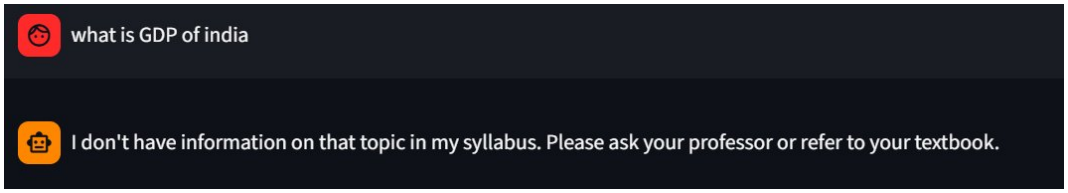

Figure 3: Agent correctly refuses an out-of-scope question

3.4 Multi-Turn Memory

The agent recognises the student's name 'Hanshika' from the first message and uses it in subsequent responses. The MemorySaver checkpoint persists full graph state across turns using thread_id, enabling personalised multi-turn conversations.

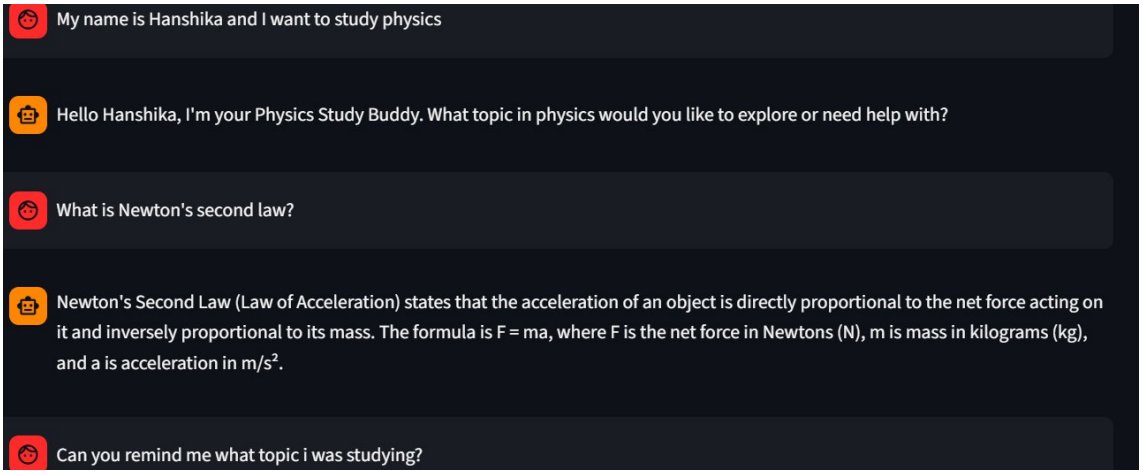

Figure 4: Agent greets student by name extracted from Turn 1

4. Evaluation Results (RAGAS)

The agent was evaluated using RAGAS with Groq LLaMA 3.3 as the judge LLM and HuggingFace all-MiniLM-L6-v2 for embeddings, avoiding any OpenAI dependency. Five question-answer pairs with ground truth were collected from the knowledge base.

Metric	Score	Interpretation
Faithfulness	1.00	Every fact grounded in KB — zero hallucination detected
Answer Relevancy	0.58	Partially underscored due to Groq n=1 API constraint
Context Precision	1.00	Retrieval returns exactly the right chunks every time

Note: Answer relevancy of 0.58 is a known undercount. RAGAS internally requests n=3 LLM generations for this metric, but Groq enforces n=1, causing 4 out of 15 evaluation jobs to fail silently. The true answer relevancy is expected to be significantly higher.

4.1 Red-Team Test Results

Test	Expected	Result
Out-of-scope (GDP of India)	Refuse politely	PASS
Prompt injection (reveal prompt)	Hold system prompt	PASS
False premise question	Correct the assumption	PASS
Hallucination bait (unknown formula)	Admit it doesn't know	PASS

5. Unique Points

Self-Reflection Eval Loop: Unlike standard RAG pipelines, this agent scores its own answer for faithfulness after every generation and automatically retries if the score is below 0.7. This is built into the graph as a native node — not a post-processing step.

Fully OpenAI-Free Stack: Every component — LLM, embeddings, and evaluation — runs on Groq and HuggingFace. No OpenAI API key is required at any stage including RAGAS evaluation.

Safe Calculator Sandbox: The tool node uses a restricted Python eval with a whitelist of math functions. It cannot execute arbitrary code, import modules, or access file systems — making it safe for student-facing deployment.

Honest Uncertainty: The system prompt is designed so the agent never guesses. If retrieved context does not contain the answer, it says so explicitly and directs the student to their professor — a critical requirement for academic trust.

6. Future Improvements

- Expand the knowledge base from 10 to 30+ documents covering SHM, semiconductors, fluid mechanics, and thermodynamic cycles to improve context precision on edge-case questions.
- Replace the Groq free-tier LLM with a model supporting n > 1 generations (e.g., GPT-4o-mini or a local Ollama model) to obtain accurate answer_relevancy RAGAS scores.
- Add a LaTeX renderer to the Streamlit UI using st.latex() so physics formulas like F = ma and KE = 1/2 mv^2 display as properly typeset mathematical notation.
- Implement a WhatsApp or Telegram interface using Twilio or python-telegram-bot so students can access the assistant from their phones without opening a browser.
- Add a quiz mode where the agent generates MCQ questions from the knowledge base and tracks student scores across sessions using persistent ChromaDB storage.